

Experiment tracking

The usual process for building an EndToEnd machine learning project involves collecting and processing raw data, analyzing it features at steps, training different algorithms, evaluating them, and deploying the best model on some platform for user access. It seems fairly straightforward, right? But in reality it is not. There are several complexities that arise along the way. Due to the circular nature of this process, its more about experimenting and trying out different things that may work

- ML is not just code. It is code plus data that we need to keep a track of. Data can be sourced from multiple storage units
- use different models and model hyperparameters
- run the same code in a different environment
- Model governance is another important aspect, where everything starting from experimentation to deployment is tracked for auditing purposes, where models are tested for speed, accuracy, drift while in production to avoid inaccuracy.

You can think of experiments as the process of building an ML model. When we say experiment run, we mean each trial in an ML experiment. So the ML experiment is actually the whole process that a data scientist may start playing with some data, models and hyperparameters. Each of these trials is an experiment run.

Experiment tracking is the process of keeping track of all the relevant information from ML experiments.

- **Organize** all the necessary components of a specific experiment. It's important to have everything in one place and know where it is so you can use them later.
- **Reproduce** past results (easily) using saved experiments.
- **Log** iterative improvements across time, data, ideas, teams, etc.

If you are working in a finance company and are tasked with creating a ml model that based on certain conditions classify if the applicant should be given a loan or not

In [1]:

```
import pandas as pd
import numpy as np
```

In [2]:

```
train_df = pd.read_csv('data.csv')
train_df.head()
```

Out[2]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_Histo
0	LP002529	Male	Yes	2	Graduate	No	6700	1750.0	230.0	300.0	1
1	LP001385	Male	No	0	Graduate	No	5316	0.0	136.0	360.0	1
2	LP001926	Male	Yes	0	Graduate	No	3704	2000.0	120.0	360.0	1
3	LP001144	Male	Yes	0	Graduate	No	5821	0.0	144.0	360.0	1
4	LP002562	Male	Yes	1	Not Graduate	No	5333	1131.0	186.0	360.0	Na

In [3]:

```
train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 430 entries, 0 to 429
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID                430 non-null    object
1   Gender                 420 non-null    object
2   Married                427 non-null    object
3   Dependents             416 non-null    object
4   Education              430 non-null    object
5   Self_Employed          410 non-null    object
6   ApplicantIncome        430 non-null    int64
7   CoapplicantIncome      430 non-null    float64
8   LoanAmount             414 non-null    float64
9   Loan_Amount_Term       422 non-null    float64
10  Credit_History          394 non-null    float64
11  Property_Area           430 non-null    object
12  Loan_Status             430 non-null    object
dtypes: float64(4), int64(1), object(8)
memory usage: 43.8+ KB
```

Binary Encoding of Categorical Variables

In [4]:

```
train_df['Gender'] = train_df['Gender'].map({'Male':0, 'Female':1})
train_df['Married'] = train_df['Married'].map({'No':0, 'Yes':1})
train_df['Loan_Status'] = train_df['Loan_Status'].map({'N':0, 'Y':1})
```

In [5]:

```
train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 430 entries, 0 to 429
Data columns (total 13 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Loan_ID               430 non-null    object
 1   Gender                420 non-null    float64
 2   Married               427 non-null    float64
 3   Dependents            416 non-null    object
 4   Education             430 non-null    object
 5   Self_Employed         410 non-null    object
 6   ApplicantIncome       430 non-null    int64
 7   CoapplicantIncome     430 non-null    float64
 8   LoanAmount            414 non-null    float64
 9   Loan_Amount_Term      422 non-null    float64
10   Credit_History         394 non-null    float64
11   Property_Area         430 non-null    object
12   Loan_Status           430 non-null    int64
dtypes: float64(6), int64(2), object(5)
memory usage: 43.8+ KB
```

Checking for Missing Values

In [6]:

```
train_df.isnull().sum()
```

Out[6]:

```
Loan_ID      0
Gender       10
Married       3
Dependents   14
Education     0
Self_Employed 20
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount   16
Loan_Amount_Term 8
Credit_History 36
Property_Area 0
Loan_Status  0
dtype: int64
```

In [7]:

```
## dropping all the missing values
train_df = train_df.dropna()
train_df.isnull().sum()
```

Out[7]:

```
Loan_ID      0
Gender       0
Married       0
Dependents    0
Education     0
Self_Employed 0
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount    0
Loan_Amount_Term 0
Credit_History 0
Property_Area 0
Loan_Status   0
dtype: int64
```

Segregating the target variable from the features

In [8]:

```
X = train_df[['Gender', 'Married', 'ApplicantIncome', 'LoanAmount', 'Credit_History']]
y = train_df.Loan_Status
X.shape, y.shape
```

Out[8]:

```
((335, 5), (335,))
```

In []:

Splitting the data

In [9]:

```
from sklearn.model_selection import train_test_split

X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=5)
```

Model Training

In [10]:

```
from sklearn.ensemble import RandomForestClassifier

model = RandomForestClassifier(max_depth=4, random_state=5)
model.fit(X_train, y_train)
```

Out[10]:

```
RandomForestClassifier(max_depth=4, random_state=5)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

Cross Validation

In [11]:

```
from sklearn.metrics import accuracy_score

pred_val = model.predict(X_val)
accuracy_score(y_val, pred_val)
```

Out[11]:

```
0.7313432835820896
```

In [12]:

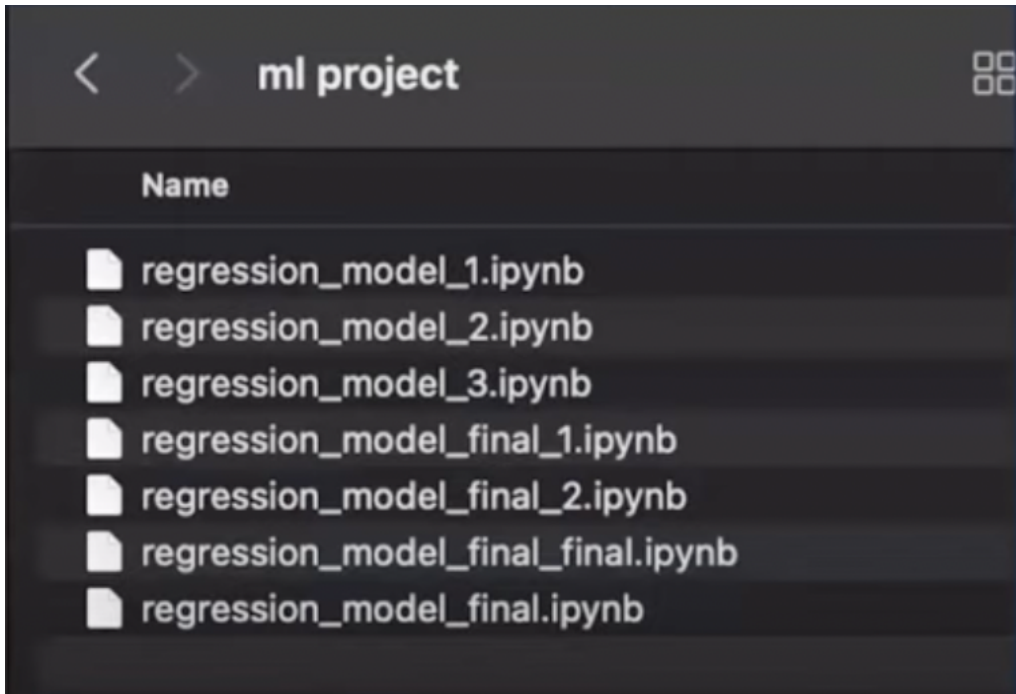
```
pred_train = model.predict(X_train)
accuracy_score(y_train, pred_train)
```

Out[12]:

```
0.8134328358208955
```

We have created a model successfully, but now we have new set of data how do you proceed on working with it

- we can change the data and run the code again
 - but we will loose the output and results from the old data
- we can create new cells below these to create a new model with the new data
 - but then when we have a lot of experiments in one file it will be really difficult finding the one we want to look at
- We can create new files for each experiment
 - but for actually comparing the results and outputs you'll still have to open each file and look into it closely



these are not the best ways of keeping track of the work and experiments that you perform, we need to create something that easy to manage, clearly shows the results and metrics, logs the changes and hyperparameters for us

ML Flow

<https://mlflow.org/> (<https://mlflow.org/>)

MLflow is an open-source platform to manage Machine Learning Lifecycle. In layman's terms, it can track and store data, parameters, and metrics to be retrieved later or displayed nicely on a web interface. Furthermore, MLflow is a framework-agnostic tool, so any ML / DL framework can quickly adapt to the ecosystem that MLflow proposes.

MLflow emerges as a platform that offers tools for tracking metrics, artifacts, and metadata.

ML flow Tracking

MLflow Tracking is an API-based tool for logging metrics, parameters, model versions, code versions, and files. MLflow Tracking is integrated with a UI for visualizing and managing artifacts, models, files, etc.

Each MLflow Tracking session is organized and managed under the **concept of runs**.

- A run refers to the execution of code where the artifact log is performed explicitly.

By default, the runs are stored in the directory where the code session is executed. However, MLflow also allows storing artifacts on a local or remote server, for better collaboration. we'll st

getting started

In [14]:

```
!pip3 install mlflow
Requirement already satisfied: databricks-cli>=0.8.7 in /Users/harshit/miniconda3/envs/dsml_env/lib/python3.9/site-packages (from mlflow) (0.17.0)
Requirement already satisfied: entrypoints in /Users/harshit/miniconda3/envs/dsml_env/lib/python3.9/site-packages (from mlflow) (0.3)
Requirement already satisfied: alembic in /Users/harshit/miniconda3/envs/dsml_env/lib/python3.9/site-packages (from mlflow) (1.8.0)
Requirement already satisfied: docker>=4.0.0 in /Users/harshit/miniconda3/envs/dsml_env/lib/python3.9/site-packages (from mlflow) (5.0.3)
Requirement already satisfied: pandas in /Users/harshit/miniconda3/envs/dsml_env/lib/python3.9/site-packages (from mlflow) (1.4.1)
Requirement already satisfied: requests>=2.17.3 in /Users/harshit/miniconda3/envs/dsml_env/lib/python3.9/site-packages (from mlflow) (2.27.1)
Requirement already satisfied: sqlalchemy in /Users/harshit/miniconda3/envs/dsml_env/lib/python3.9/site-packages (from mlflow) (1.4.39)
Requirement already satisfied: pyjwt>=1.7.0 in /Users/harshit/miniconda3/envs/dsml_env/lib/python3.9/site-packages (from databricks-cli>=0.8.7->mlflow) (2.4.0)
Requirement already satisfied: oauthlib>=3.1.0 in /Users/harshit/miniconda3/envs/dsml_env/lib/python3.9/site-packages (from databricks-cli>=0.8.7->mlflow) (3.2.0)
Requirement already satisfied: six>=1.10.0 in /Users/harshit/miniconda3/envs/dsml_env/lib/python3.9/site-packages (from databricks-cli>=0.8.7->mlflow) (1.15.0)
```

In [15]:

```
import mlflow
import os
```

we'll start by setting up our experiment name under which we wanna perform all our work

- An MLflow experiment is the primary unit of organization and access control for MLflow runs; all MLflow runs belong to an experiment. Experiment:{run,run.....run}
- Experiments let you visualize, search for, and compare runs, as well as download run artifacts and metadata for analysis in other tools.
- An MLflow run corresponds to a single execution of model code. Each run records the some information about that particular trial:

In [16]:

```
mlflow.set_experiment("loan_status")
```

Out[16]:

```
<Experiment: artifact_location='file:///Users/abdulhad_scaler/jupyter/mlruns/1', experiment_id='1', lifecycle_stage='active', name='loan_status', tags={}>
```

Next, you can start to think about what do you want to keep track in your analysis/experiment. MLflow categorizes these into:

- **Parameters** (via `mlflow.log_param()`). Parameters are variables that you change or tweak when tuning your model.
- **Metrics** (using `mlflow.log_metric()`). Metrics are values that you want to measure as a result of tweaking your parameters. Typical metrics that are tracked can be items like F1 score, RMSE, MAE etc.
- **Artifacts** (using `mlflow.log_artifact()`). Artifacts are any other items that you wish to store. Typical artifacts to keep track of are PNGs of graphs, plots, confusion matrix, and also pickled model files

Params are something you want to tune based on the metrics, whereas tags are some extra information that doesn't necessarily associate with the model's performance. there's no hard constraint on which to use to log which; they can be used interchangeably without error.

In [17]:

```
with mlflow.start_run():
    model_rf = RandomForestClassifier(max_depth=4, random_state=5)
    model_rf.fit(X_train, y_train)

    pred_val = model_rf.predict(X_val)
    val_acc=accuracy_score(y_val, pred_val)

    pred_train = model_rf.predict(X_train)
    train_acc=accuracy_score(y_train, pred_train)

    mlflow.set_tag('mlflow.runName', 'first_run')
    mlflow.log_param('max_depth', 4)
    mlflow.log_metric('val_acc', val_acc)
    mlflow.log_metric('train_acc', train_acc)

    mlflow.sklearn.log_model(model_rf, "model")
```

```
-----
PermissionError                                Traceback (most recent call last)
/var/folders/v3/9qnnmcxd5rdbhy0swt0r_1lm0000gn/T/ipykernel_43059/595822507.py in <module>
     14     mlflow.log_metric('train_acc', train_acc)
     15
--> 16     mlflow.sklearn.log_model(model_rf, "model")
     17
     18

~/miniconda3/envs/dsml_env/lib/python3.9/site-packages/mlflow/sklearn/_init_.py in log_model(sk_model, artifact_path, conda_env, code_paths, serialization_format, registered_model_name, signature, input_example, await_registration_for, pip_requirements, extra_pip_requirements)
     391     mlflow.sklearn.log_model(sk_model, "sk_models")
     392     """
--> 393     return Model.log(
     394         artifact_path=artifact_path,
     395         flavor=mlflow.sklearn,
```

we can use this with command to start the ml flow run and whatever we do inside of that start_run indent will be tracked

inside that we create our first model and log the different parameters and metric for that model we set the name of the run and log the max depth of the rf model and also the acc score. All of the parameters and models are stored in files in the experiment folder with each runs having seperate folders. you can open those files to see the stored data

```
jupyter — Python < mlflow ui — 86x24
abdulahad_scaler@Abduls-Air jupyter % mlflow ui
[2022-10-13 01:28:33 +0530] [61216] [INFO] Starting gunicorn 20.1.0
[2022-10-13 01:28:33 +0530] [61216] [INFO] Listening at: http://127.0.0.1:5000 (61216)
[2022-10-13 01:28:33 +0530] [61216] [INFO] Using worker: sync
[2022-10-13 01:28:33 +0530] [61217] [INFO] Booting worker with pid: 61217
[2022-10-13 01:28:33 +0530] [61217] [INFO] Handling signal: winch
```

mlflow ui

MLflow also provides the option to view all the runs and experiments on a web based ui that is really easy to use and see the logged data. Launch the MLflow tracking UI for local viewing of run results. In the folder where you have this experiments run the command **mlflow ui** this will start an ml flow ui server that is by default open at port 5000 on your localhost or 127.0.0.1 you can change the port by using -p port_num along with the command eg: **mlflow ui -p 8899**

- open the correct link or copy the provided url from the command

The screenshot shows the MLflow 1.29.0 Experiments web UI. The 'loan_status' experiment is selected. The interface includes a search bar, a list of experiments (Default and loan_status), and a table of runs. The table shows one matching run named 'first_run' created 8 minutes ago by user 'abdulahad_...'. The table columns include Created, Duration, Run Name, User, Source, Version, Models, and Metrics. A search bar and filter options are visible at the top.

Created	Duration	Run Name	User	Source	Version	Models	Metrics
8 minutes ago	2.2s	first_run	abdulahad_...	ipykerne...	-	sklearn	0.81:


here is the web ui launched on a browser, as you can see we are under the loan_status experimnet name and have a run that we created with the name first_run. there are several other informations as welllike the source code that we used, the user thatcreated that run and the model that we stored

loan_status > first_run

first_run

Run ID: 7d52e1ff68a34342b25ed2da8e0b69df

Date: 2022-10-13 01:20:18

Source:  ipykernel_launcher.py

User: abdulahad_scaler

Duration: 2.2s

Status: FINISHED

Lifecycle Stage: active


> Description [Edit](#)


> Parameters (1)


> Metrics (2)


> Tags


> Artifacts


▼  model

 MLmodel

 conda.yaml

 model.pkl

 python_env.yaml

 requirements.txt

Full Path:file:///Users/abdulahad_scaler/jupyter/mlruns/1/7d52e1ff68a34342b25ed2da8e0b69df/artifacts/...



Register Model

if we click on any particular run we can see more details about that run. we have here the all the details that we logged for that particular model in a very easy to understand fashion

▼ Parameters (1)

Name	Value
max_depth	4

▼ Metrics (2)

Name	Value
train_acc 	0.813
val_acc 	0.731

Load the new data and proceed further

localhost:8888/notebooks/Experiment tracking.ipynb

7/11

```
In [53]:

train_df = pd.read_csv('data_new.csv')
train_df.head()

train_df['Gender'] = train_df['Gender'].map({'Male':0, 'Female':1})
train_df['Married'] = train_df['Married'].map({'No':0, 'Yes':1})
train_df['Loan_Status'] = train_df['Loan_Status'].map({'N':0, 'Y':1})

train_df = train_df.dropna()

X = train_df[['Gender', 'Married', 'ApplicantIncome', 'LoanAmount', 'Credit_History']]
y = train_df.Loan_Status
X.shape, y.shape

from sklearn.model_selection import train_test_split

X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=5)
```

```
In [54]:

with mlflow.start_run():
    model_rf = RandomForestClassifier(max_depth=4, random_state=5)
    model_rf.fit(X_train, y_train)

    pred_val = model_rf.predict(X_val)
    val_acc=accuracy_score(y_val, pred_val)

    pred_train = model_rf.predict(X_train)
    train_acc=accuracy_score(y_train, pred_train)

    mlflow.set_tag('mlflow.runName', 'new_data')
    mlflow.log_param('max_depth',4)
    mlflow.log_metric('val_acc',val_acc)
    mlflow.log_metric('train_acc',train_acc)
    mlflow.set_tag('data file','data_new.csv')


    mlflow.sklearn.log_model(model_rf, "model")
```

> Description [Edit](#)

↻ Refresh

Compare

Delete



 Download CSV

↓ Created ▾


All time ▾

Columns ▾

Only show differences ☐

  metrics.rmse < 1 and params.model = "tree"

Search

 Filter

Clear

Showing 2 matching runs

					Metrics		Parameters	Tags
<input type="checkbox"/>	↓ Created	Duration	Run Name	s	train_acc	val_acc	max_depth	data file
<input type="checkbox"/>	🕒 6 seconds ago	1.8s	new_data	arn	0.839	0.792	4	data_new.c...
<input type="checkbox"/>	🕒 1 hour ago	2.2s	first_run	arn	0.813	0.731	4	-

Load more

if we go back to the web ui we can see that we have another run logged with the information we have we changed we added a new name and the name of the datafile

Now if we want to tune the RF model

In [61]:

```
def mlflow_runs(n_est,max_dep,i):
    with mlflow.start_run():

        model_rf = RandomForestClassifier(n_estimators=n_est, max_depth=max_dep, random_state=5)
        model_rf.fit(X_train, y_train)

        pred_val = model_rf.predict(X_val)
        val_acc=accuracy_score(y_val, pred_val)

        pred_train = model_rf.predict(X_train)
        train_acc=accuracy_score(y_train, pred_train)

        run="hyperparameter_run_"+str(i)
        mlflow.set_tag('mlflow.runName',run)
        mlflow.log_param('n_estimators',n_est)
        mlflow.log_param('max_depth',max_dep)
        mlflow.log_metric('val_acc',val_acc)
        mlflow.log_metric('train_acc',train_acc)
        mlflow.set_tag('data file','data_new.csv')

        mlflow.sklearn.log_model(model_rf, "model")
```

In [62]:

```
mlflow_runs(10,2,1)
mlflow_runs(20,2,2)
mlflow_runs(40,2,3)
mlflow_runs(10,4,4)
mlflow_runs(20,4,5)
mlflow_runs(40,4,6)
mlflow_runs(10,8,7)
mlflow_runs(20,8,8)
mlflow_runs(40,8,9)
```

Showing 13 matching runs

				Metrics		Parameters		Tags
<input type="checkbox"/>	↓ Created	Duration	Run Name	train_acc	val_acc	max_depth	n_estimators	data file
<input type="checkbox"/>	✔ 31 seconds ago	0.8s	hyperpara...	0.914	0.802	8	40	data_new.c...
<input type="checkbox"/>	✔ 32 seconds ago	0.8s	hyperpara...	0.914	0.771	8	20	data_new.c...
<input type="checkbox"/>	✔ 33 seconds ago	0.8s	hyperpara...	0.909	0.75	8	10	data_new.c...
<input type="checkbox"/>	✔ 33 seconds ago	0.8s	hyperpara...	0.833	0.792	4	40	data_new.c...
<input type="checkbox"/>	✔ 34 seconds ago	0.8s	hyperpara...	0.833	0.792	4	20	data_new.c...
<input type="checkbox"/>	✔ 35 seconds ago	0.7s	hyperpara...	0.828	0.792	4	10	data_new.c...
<input type="checkbox"/>	✔ 36 seconds ago	0.8s	hyperpara...	0.807	0.813	2	40	data_new.c...
<input type="checkbox"/>	✔ 36 seconds ago	0.7s	hyperpara...	0.807	0.813	2	20	data_new.c...
<input type="checkbox"/>	✔ 38 seconds ago	1.7s	hyperpara...	0.807	0.813	2	10	data_new.c...
<input type="checkbox"/>	✖ 2 minutes ago	267ms	hyperpara...	-	-	-	-	-
<input type="checkbox"/>	✖ 3 minutes ago	105ms	lyrical-jay-...	-	-	-	-	-
<input type="checkbox"/>	✔ 38 minutes ago	1.8s	new_data	0.839	0.792	4	-	data_new.c...
<input type="checkbox"/>	✔ 2 hours ago	2.2s	first_run	0.813	0.731	4	-	-

Load more

here we can see there are 9 new runs that show how our model performed

- we can see that increasing the number of tress improves the model a lot
- if we have a deep model with less number of trees it seems to overfit because the train accuracy is very high but the valaccuracy is low you can also see there are two failed runs so they have no data associated with them

now if you want to try out another model like knn for this task

In [77]:

```
from sklearn.neighbors import KNeighborsClassifier
import matplotlib.pyplot as plt
from sklearn.metrics import ConfusionMatrixDisplay

with mlflow.start_run():
    knn_model= KNeighborsClassifier(n_neighbors=5)
    knn_model.fit(X_train, y_train)

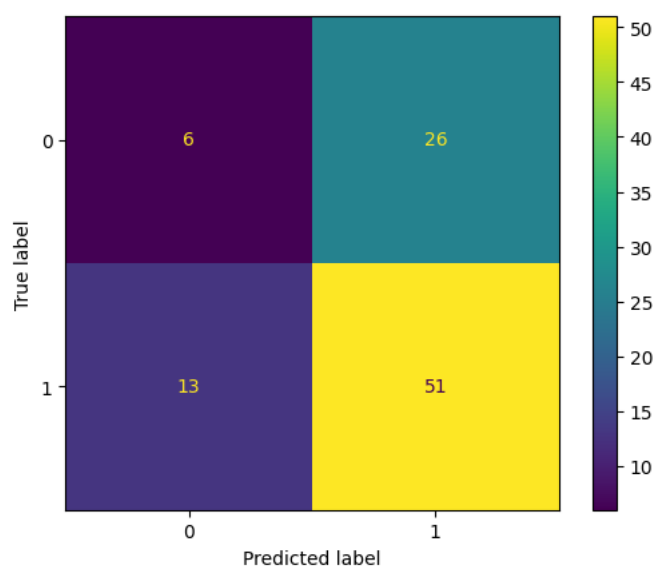
    pred_val = knn_model.predict(X_val)
    val_acc=accuracy_score(y_val, pred_val)

    pred_train = knn_model.predict(X_train)
    train_acc=accuracy_score(y_train, pred_train)

    run="KNN"
    mlflow.set_tag('mlflow.runName',run)
    mlflow.log_param('neighbors',5)
    mlflow.log_metric('val_acc',val_acc)
    mlflow.log_metric('train_acc',train_acc)
    mlflow.set_tag('data file','data_new.csv')

    cm=ConfusionMatrixDisplay.from_predictions( y_val,pred_val)
    cm.figure_.savefig('confusion_mat.png')
    mlflow.log_artifact('confusion_mat.png')

    mlflow.sklearn.log_model(knn_model, "model")
```



Run ID: b97f4fab7a649bf998547c893207d6a

Date: 2022-10-13 04:56:58

User: abduhad_scaler

Duration: 1.7s



Lifecycle Stage: active

> Description [Edit](#)

> Parameters (1)

Name	Value
neighbors	5

> Metrics (2)

Name	Value
train_acc 	0.766
val_acc 	0.594

after all the testing and trying we can say that we will chose the random forest model with max depth =8 and number of trees=40